AAFI MALEK

Applied AI Engineer

aafimalek2032@gmail.com | +91 7574067294 | github.com/Aafimalek | linkedin.com/in/aafi-malek

Professional Summary

Applied AI Engineer building Python-based LLM applications, multi-agent workflows, enterprise RAG/GraphRAG, and real-time voice AI. Experienced with FastAPI, LangGraph, MCP, Google ADK, AWS Bedrock, secure tool execution, and evaluation-driven deployment. Delivers end-to-end AI features with observable performance, reliable APIs, controlled data access, and automated testing.

Technical Skills

LLMs & Agentic AI: GPT/OpenAI, Claude/Anthropic, Gemini, open-weight LLMs, LangGraph, LangChain, Google ADK, AWS Strands, MCP, tool/function calling, memory, human-in-the-loop, structured output, prompt/context engineering

RAG & Document AI: RAG, GraphRAG, hybrid search, embeddings, semantic retrieval, chunking, document extraction, source citations, LandingAI ADE, Bedrock Knowledge Bases, Neo4j, Pinecone, Qdrant, OpenSearch

Backend & Data: Python, FastAPI, Flask, REST APIs, asyncio, WebSockets, SSE, webhooks, PostgreSQL, Neon, MongoDB, Redis, Pandas, NumPy, Pydantic

Cloud & DevOps: AWS Bedrock/AgentCore, Lambda, API Gateway, SQS, EventBridge, S3, IAM, CloudWatch, ECS/Fargate, Docker, Git, GitHub Actions CI/CD

Evaluation & Voice AI: Langfuse, LangSmith, golden datasets, LLM-as-a-judge, regression evaluation, tool-call validation, latency/cost tracing, Vapi, LiveKit, Twilio, Deepgram, n8n

Professional Experience

AI/ML Engineer (Full-time) | *DatumQuest Private Limited* **Jul 2026 – Present**

- Owned Python/FastAPI backends for 4 applied AI initiatives across agent automation, voice operations, enterprise knowledge retrieval, and multimodal content; integrated REST APIs, async jobs, and PostgreSQL.
- Built durable LangGraph/Google ADK agents with MCP tool discovery, schema-validated actions, memory, recovery, permission scopes, and human approval; improved task completion by 22% on a 120-scenario internal evaluation set.
- Designed hybrid GraphRAG with Neo4j relationships and Pinecone/Qdrant search; added tenant-aware retrieval, evidence citations, and grounded-answer checks, reducing unsupported responses by 30% in test scenarios.
- Integrated Vapi/Google ADK and Twilio/Deepgram streaming with AWS Lambda, SQS, EventBridge, Google Calendar, and n8n, enabling real-time voice workflows and automated post-call actions.
- Established observability and quality controls with Langfuse traces, CloudWatch logs, golden datasets, and GitHub Actions CI/CD; shortened release-validation time by 35% through repeatable regression tests.
- Collaborated with product stakeholders on API contracts, deployment readiness, acceptance criteria, technical documentation, and operational runbooks.

AI Engineer Intern | *DatumQuest Private Limited* **Jan 2026 – Jun 2026**

- Developed 12+ FastAPI endpoints, webhook integrations, async tasks, and PostgreSQL/Redis workflows; added request validation, API documentation, and unit/integration tests.
- Built PDF ingestion pipelines covering extraction, chunking, embeddings, vector indexing, and multilingual document Q&A; evaluated retrieval quality using a 100-query reference set.
- Integrated live voice-agent proof-of-concepts with external speech APIs and webhook callbacks; added error handling, structured logs, and automated test cases.
- Assisted with Docker-based delivery, Git workflows, API specifications, and debugging of asynchronous services through staging and release testing.

Projects

AI Voice Agent Operations & Observability Platform *Google ADK, Vapi, AWS*

- Built a unified operations platform for real-time AI voice and chat agents, connecting FastAPI services, Google ADK and Vapi runtimes, and a PostgreSQL-backed management dashboard.
- Implemented role-based access control (RBAC), agent configuration APIs, conversation records, and operational workflows supporting managed Vapi and custom Google ADK agents.
- Integrated Twilio Media Streams and Deepgram STT/TTS for live speech interactions, including call events, webhook handling, and downstream calendar scheduling through Google Calendar and n8n.
- Implemented asynchronous event processing with AWS Lambda, API Gateway, and SQS to coordinate call automation, post-call tasks, and external-service integrations.
- Added hybrid GraphRAG using Neo4j relationships and Pinecone vector retrieval to answer questions involving brokers, groups, agents, listings, and related business records.
- Instrumented Langfuse traces, retrieval evidence, and tool execution for debugging and conversation-quality reviews; reduced manual post-call follow-ups by 40% in pilot workflows.

Content Generation & Retrieval Harness *LandingAI, Qdrant, LangGraph*

- Developed a multimodal AI application that ingests campaign-specific websites, PDFs, images, and brand assets and transforms them into searchable source material for grounded content generation.
- Implemented asynchronous FastAPI workflows for document extraction, indexing, retrieval, evaluation, and media jobs, keeping ingestion separate from generation requests.
- Used LandingAI Agentic Document Extraction and website capture to extract structured text, factual claims, visuals, and source metadata from heterogeneous business documents.
- Built semantic retrieval with OpenAI embeddings and Qdrant; maintained campaign records, extracted facts, citation provenance, generation jobs, and evaluation results in Neon PostgreSQL.
- Stored original documents, extracted assets, and generated content in Amazon S3; orchestrated LangGraph agents for source selection, generation, formatting, and unsupported-claim checks.
- Added output checks for source misuse and tone/brand consistency across LinkedIn, X, and Instagram content; reduced manual research and preparation effort by 45% in pilot tests.

Education

Bachelor of Engineering, Artificial Intelligence & Machine Learning *L. D. College of Engineering*